

Wrapping up the Pipeline

All the nitty gritty details you wished to get rid of...


Elmar Eisemann

Delft University of Technology (TU Delft)

- A long time ago in a galaxy far far away...


Pinhole camera

- Box with hole
- Perfect image for “point-sized” hole




Graphics Pipeline

- Highly **parallelizable** → **GPUs**




Triangle Mesh

- Models are typically lists of **triangles**



Simplified Graphics Pipeline

- **Projection:** Transform coordinates to screen



Summary

- First steps into the Graphics Pipeline
 - Homogenous Coordinates
 - Matrix Stacks
 - Camera Model

Complete Camera Model

- Finally:

$$P = \underbrace{\begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix}}_{\text{image}} \underbrace{\begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix}}_{\text{projection}} \underbrace{\begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \\ 0 & 0 & 0 & 1 \end{pmatrix}}_{\text{orientation/location}}$$

$$\begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

Complete Camera Model

- Finally:

$$P = \begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

$$\begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \end{pmatrix}$$


Complete Camera Model

- Finally (notation often used in vision literature):

$$P = \begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \end{pmatrix}$$

intrinsic / internal parameters extrinsic / external parameters

Change settings Move camera



Questions?

IF TRAIN A LEAVES BOSTON AT 9:27 PM
HEADING WEST AT 173 MPH AND TRAIN B
LEAVES MILWAUKEE AT 10:38 AM HEADING
EAST AT 123 MPH WITH A STEADY
NORTH WIND BLOWING AT 17 MPH,
HOW LONG WILL IT TAKE YOU
TO FIND ANOTHER JOB?



Another layoff at the textbook publishing company.

Copyright 2001 by Randy Glasbergen. www.glasbergen.com

It could have been so simple...


- What is the problem of this matrix for the Graphics Pipeline?

$$P = \begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \end{pmatrix}$$

We cannot eliminate depth!

- We need to keep a Z!


Compare new distance to stored distance
Update pixel only if new distance is nearer



- Do we need the “true depth”?
- Any monotonic mapping would work!

Problem:

- A 3D scene is infinite...
- How do we represent Z?
- Solution add a near and far clipping plane!




Camera in Matrix Representation

- Camera space:

screen mapping \rightarrow $\begin{bmatrix} f & 0 & 0 & 0 \\ \text{aspect} & f & 0 & 0 \\ 0 & 0 & \text{near} + \text{far} & 2\text{near}\text{far} \\ 0 & 0 & \text{near} - \text{far} & \text{near} - \text{far} \end{bmatrix}$ \leftarrow Z mapping

In practice orientation along z-axis is inverted
...don't bother too much...




Camera in Matrix Representation

- What does this matrix do?

– Vertices multiplied with this matrix change their coordinates and are projected onto the screen.

– In other words, the entire scene is transformed!



Camera Space

- Perspective Frustum is mapped on a cube





Illustration of the Camera Mapping



$$\begin{bmatrix} f & 0 & 0 & 0 \\ \text{aspect} & f & 0 & 0 \\ 0 & 0 & \text{near} + \text{far} & 2\text{near}\text{far} \\ 0 & 0 & \text{near} - \text{far} & \text{near} - \text{far} \end{bmatrix}$$

- Verify:
- 1) (0,0,-near,1) has depth -1, and (0,0,-far,1) has depth 1
- 2) What does (0,0,-(far+near)/2,1) map to?
- 3) Where are those points in the image?

You will notice, the Z-Buffer is not linear, the zero depth will not map to the frustum center

Graphics Pipeline Camera

Historically the graphics pipeline has a projection and a model/view matrix and a viewport matrix to map to the pixels.

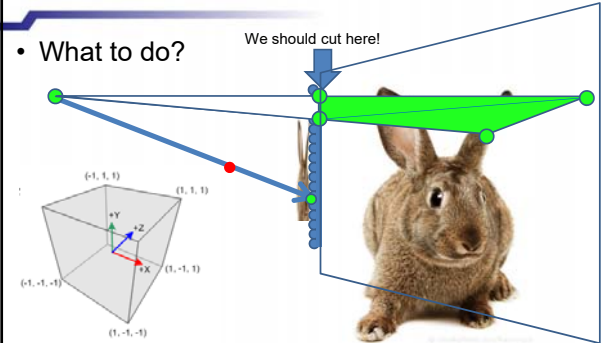
$$\begin{bmatrix} a_x & 0 & 0 & x_0 \\ 0 & a_y & 0 & y_0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} * \begin{bmatrix} \frac{f}{\text{aspect}} & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & \frac{\text{near} + \text{far}}{\text{near} - \text{far}} & \frac{2\text{near far}}{\text{near} - \text{far}} \\ 0 & 0 & -1 & 0 \end{bmatrix} * \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

First the model/view matrix is applied, then the projection matrix, then the viewport

What is what?

Points behind camera?


- What to do?



- Test if triangle lies in the cube!


Clip Space

- Create clipped triangles



Simplified Graphics Pipeline

- Rasterization: Fill screen pixels



+ Depth Test

Summary

- Graphics Pipeline
 - Homogenous Coordinates
 - Matrix Stacks
 - Camera Model
 - Clipping/Rasterization
- Standard Pipeline easy to use!

Images are not very exciting yet...



Shading

How to transform




in



?

Questions?



Graphics API

- Two major types: DirectX & OpenGL
 - Allows us to send commands to the card
 - Relatively close to hardware


Main difference: OpenGL is a state machine
-> activate something and it stays activated

- We will use OpenGL


Making beautiful pictures...

TODAY'S QUESTION:

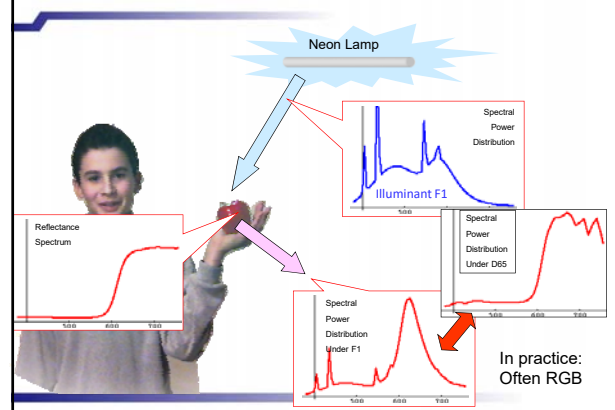
- Given point properties (i.e., position, normal, and other attributes)
- How to compute a realistic color???



What is Color?




What is Color?




Reflection

- What happens when the light hits a surface?




Reflection

- What happens when the light hits a surface?




Reflection

- What happens when the light hits a surface?




Reflection

- What happens when the light hits a surface?



Reflection


- What happens when the light hits a surface?



Today:
Only interested in one
particular direction!

Making beautiful pictures...

- Given point properties (i.e., position, normal, and other attributes)
- Apply shading model



Acquired BRDFs

- Big databases

Often costly, much data,
hard to use for artists
e.g., "Can you make the blue darker?"



Simplified Models

- Mathematically describe Material Properties
- Sum of 3 terms
 - Ambient
 - Diffuse
 - Specular
 - Phong
 - Blinn-Phong

Color - Recap

- Remember:
 - Visual system uses 3 cone types for color
- Each component can be treated separately.
Hence, in the following, we use a single scalar.
(Imagine it to be red, green, or blue...)

Ambient Term


- Represent "scene light"
 - Skylight
 - Light reflected from neighboring surfaces
- Model:
 - Extremely simple

Componentwise multiplication (RGB)! Each component can be treated separately!

$$A = I_a K_a$$

Light property (RGB) →

Surface property (RGB) →



Ambient Term

- Very simplistic
 - No real physical basis...
 - No indications on the shape of an objet!



- But often used in practice
 - Emulates indirect light from scene
 - Elements in shadow are not completely black


Diffuse Term

- Lambert Surfaces – isotropic reflection
 - Reflection uniformly in all directions
 - Conserves energy
- Model:
 - Uses local surface orientation

Light property (RGB) →


$$D = I_d K_d \cos \theta$$

Surface property (RGB) →



Diffuse Term

- Hypothesis: isotropic a.k.a lambertian
- Shading varies along surface
 - Gives information about shape




- Does not depend on observer position!

Diffuse Term

- Where does the cosine come from?

$$D = I_d K_d \cos \theta$$



- Imagine light direction as a set of parallel rays
- What do you observe when tilting the rays on a small surface patch?




Specular Term

- Represent glossy surfaces
 - Ideal case: mirror
- Model:
 - reflection around mirror ray
 - Exponential decrease

$$S(\phi) = I_s K_s (\cos \phi)^n$$



Specular Term




Example

$$A = I_a K_a \quad D = I_d K_d \cos \theta$$


- K_d diffuse: (0.9,0,0)
- I_d diffuse light: (0.9,0.5,1.0)
- K_a ambiante: (0,0,0.1)
- I_a ambiante light: (1,1,0.1)

- Normal at x: (0,0,1)
- Position x: (0,0,0)
- light position : (0,0,10)
- Resulting color: (0.81,0,0.01)




The other extreme...

- Non-photorealistic materials



- Cel-shading: threshold on the diffuse shading

Questions?



And now in practice?

- We know how to compute shading of a point, but how is it applied on a mesh?

Shading

- Early days - compute color per face:
Flat shading produces "facets"




- Later – compute color per vertex: produces
Gouraud Shading produces a smooth look



Phong shading

- Today: compute result per pixel
- Phong Shading leads to smooth specularities



Diccan.com

On the practical side: Shading types

- How are the three different types computed?

– Flat shading

- Applies Phong Model to produce a color **per face**

– Gouraud shading – default in OpenGL

- Applies Phong to produce a color **per vertex**
- Interpolate color from vertices over triangle


– Phong shading

2 MEANINGS!!!

- Interpolate parameters of Phong model :


- Parameters are :
 - Normal
 - Vector to the lightsource (light vector)

- Applies Phong to produce a color **per pixel**




Per vertex Gouraud vs. Phong shading Per pixel

- Phong more expensive than Gouraud
 - Usually: more pixels than vertices
- Phong more beautiful and standard today
 - Captures specularities between faces



Diccan.com



Questions?




Something is still missing...

Misses quick material changes




Textures

- Can be used to provide, e.g.,
 - parameters for ambient/diffuse material models




Textures

- Can be used to provide, e.g.,
 - Colors for ambient/diffuse
 - Normals (*bump mapping*)




Textures

- Can be used to provide, e.g.,
 - Colors for ambient/diffuse
 - Normals (*bump mapping*)



Textures


- Image mapped on the surface via texture coordinates
 - Specified at each vertex `glTexCoord(123){fi}`
 - Interpolated over triangles



```
TextureCoords(0.2f, 0.3f);
Vertex(-1.0f, 0.0f);


TextureCoords(0.5f, 0.4f);
Vertex(+1.0f, 0.0f);

TextureCoords(0.1f, 0.8f);
Vertex( 0.0f, 1.0f);
```




How to define Texture Coordinates?

- We will provide tex coords by hand!



Texture coordinates are handled like attributes and interpolated over triangle!!!



Textures

- Extremely useful and efficient!
- What are the problems?
- We will see about this next time...

Questions?



Summary

- Overview of the “practical” graphics pipeline
- Graphics Pipeline
 - Clipping of elements outside the frustum
 - Shading
 - material models (Diffuse, Specular)
 - Interpolation of shading results (Flat, Gouraud, Phong)
- Textures
 - Applications
 - Texture coordinates

Questions?

